

IMPROVEMENT OF EPSILON COMPLEXITY ESTIMATION AND AN
APPLICATION TO SEIZURE PREDICTION

A thesis presented to the faculty of
San Francisco State University
In partial fulfilment of
The Requirements for
The Degree

Master of Arts
In
Mathematics

by

Nathanael Aff

San Francisco, California

August 2017

CERTIFICATION OF APPROVAL

I certify that I have read *IMPROVEMENT OF EPSILON COMPLEXITY ESTIMATION AND AN APPLICATION TO SEIZURE PREDICTION* by Nathanael Aff and that in my opinion this work meets the criteria for approving a thesis submitted in partial fulfillment of the requirements for the degree: Master of Arts in Mathematics at San Francisco State University.

Dr. Alexandra Piryatinska
Associate Professor of Mathematics

Dr. Tao He
Assistant Professor of Mathematics

Dr. Chun Kit Lai
Assistant Professor of Mathematics

IMPROVEMENT OF EPSILON COMPLEXITY ESTIMATION AND AN APPLICATION TO SEIZURE PREDICTION

Nathanael Aff
San Francisco State University
2017

The ε -complexity of a continuous function is a measure of the amount of information needed to reconstruct a function with an absolute error not larger than ε . For Hölder class functions, ε -complexity is characterized by a pair of real numbers, the complexity coefficients. The complexity coefficients have been shown to be useful features for the segmentation and classification of time series. In this work, we extend the set of approximation methods used in estimating the complexity coefficients. The performance of these approximation methods is tested on a number of simulated time series. In addition, we test the conjecture that, for a given generating mechanism, the mean of the complexity coefficients is constant. For our set of simulations, we find that the mean of the estimated complexity coefficients is constant on a constant Hölder class of functions. Finally, we apply the ε -complexity coefficients to the prediction of seizures in epileptic mice. We use this technique to identify which EEG signal preceded a seizure with over 80% accuracy.

I certify that the Abstract is a correct representation of the content of this thesis.

Chair, Thesis Committee

Date

ACKNOWLEDGMENTS

I thank my advisor Dr. Alexandra Piryatinska her for her guidance and my two committee members, Dr. Tao He and Dr. Chun-Kit Lai, I thank for their patience.

TABLE OF CONTENTS

1	Introduction	1
2	Background	5
2.1	Epsilon-Complexity	5
2.2	Time Series	10
2.3	Fractal Dimension	14
2.4	Simulations	16
2.5	Approximation Methods	20
3	Approximation Methods	23
3.1	Approximation and Classification Accuracy	24
3.2	Computational Efficiency	30
4	Interpreting the Complexity Coefficients	33
4.1	The Complexity Coefficients	34
4.2	Hölder Class and Fractal Dimension	37
5	Seizure Prediction	45
5.1	Introduction	46
5.2	EEG Data	47
5.3	EEG Features	48
5.4	Classification Procedure	49
5.5	Model Performance	52
5.6	Feature Importance	56
	Bibliography	62

LIST OF TABLES

Table	Page
3.1 Initial parameters for each simulation group.	25
3.2 Mean total approximation error for 30 samples of each simulation. . .	29
3.3 Classification errors using complexity coefficients A and B for each approximation method.	29
4.1 Complexity coefficients and fractal dimension estimates for linear and sinusoidal functions.	35
4.2 Complexity coefficients and fractal dimension estimates for Gaussian noise.	36
5.1 Classification performance of baseline classifier.	53
5.2 Mean and standard deviation of balanced accuracy for all models. . .	55
5.3 Mean and standard deviation of specificity for all models.	55
5.4 Results of an unpaired Wilcoxon test.	59
5.5 Normalized difference of medians (Seizure – Non-seizure) for each feature and trial.	60

LIST OF FIGURES

Figure	Page
3.1 Sample functions from each simulation group.	26
3.2 Linear regression plots for the lifting scheme.	27
3.3 The two simulation groups plotted in the complexity coefficient space.	28
3.4 Computational time of each approximatio method as a function of the size of the log of the input.	31
4.1 Simple linear and sinusoidal functions (left) and with added noise (right).	34
4.2 The log-log linear regression of approximation errors against sample proportion S_h for simple linear and sinusoidal functions.	35
4.3 The Weierstrass and random-phase Weierstrass function for $\alpha =$ $\{0.20, 0.42, 0.80\}$	38
4.4 Complexity coefficients and fractal dimension for values of the Weier- strass α parameter.	39
4.5 Complexity coefficients and fractal dimension for values of the random- phase Weierstrass α parameters.	39
4.6 Complexity coefficients and fractal dimension for values of the random- phase Weierstrass α parameters.	40
4.7 Fractional Brownian motion and the Cauchy process with $\alpha = \{0.20, 0.42, 0.80\}$ for a constant β parameter.	41
4.8 Complexity coefficients and fractal dimension for values of the α pa- rameter of fractional Brownian motion.	41
4.9 Complexity coefficients and fractal dimension for various α parame- ters for a constant β parameter of the Cauchy process.	42

4.10	Mean value of the complexity coefficients and fractal dimension for each α value of the Cauchy process.	43
4.11	The mean value of the complexity coefficients and fractal dimension for each β value of the Cauchy process.	44
5.1	Spectral features segmented on changes in the B complexity coefficient.	52
5.2	Diagnostic values with 95% bootstrap confidence intervals for each classifier trained on channel 1 features.	54
5.3	Best ROC curves for classifiers B and 8. On the right is the smoothed curve.	56
5.4	Normalized variable importance for each partition method and channel.	57
5.5	Feature distribution for channels 1(top) and 3(bottom).	58

Chapter 1

Introduction

Many natural phenomena generate time series that exhibit complex functional and statistical characteristics. For example, an electroencephalogram (EEG) records the electrical potential generated by the synchronized firing of neurons. EEG signals are marked by both transient variations in waveforms and regime changes —distinct breaks in the character and statistical distribution of the EEG. While the exact generating mechanisms of variations in EEG may not be known, characteristics of the observed signal can be used to identify changes in the underlying dynamics.

The theory of ε –complexity as developed by Darkhovsky and Piryatinska provides a method of identifying regime changes in complex signals such as EEG. The definition of ε –complexity is related to Kolmogorov’s definition of the complexity of a sequence. Roughly, Kolmogorov identified the complexity of a discrete sequence as the size of the program, or the amount of information needed to produce that sequence. The ε –complexity of a continuous function, on the other hand, is the amount of information needed to reconstruct that function with an absolute error not greater than ε . In the latter case, we consider the proportion of the function needed for reconstruction analogous to information in Kolmogorov’s theory. Dark-

Darkhovsky and Piryatinska have shown that the ε -complexity of a continuous function can be identified by two real numbers[5]. These are the ε -complexity coefficients, or simply the complexity coefficients of the function. In practical applications, a signal like EEG is given by a some finite set of samples. The theory of ε -complexity is adapted to a discrete setting by considering discrete sequences, such as time series, as restrictions of continuous functions to a uniform grid.

The procedure for estimating the ε -complexity of a function entails iteratively approximating that function on a more sparsely sampled set of points at each iteration. The initial implementation of the procedure used piecewise polynomials to approximate the original function. We implement the estimation procedure with an enlarged set of approximation methods —basis splines, cubic splines, and an interpolating subdivision method termed the lifting scheme. We test the estimation of ε -complexity on simulated data using each approximation method. First, we compare the reconstruction accuracy of each method. Then we test the performance of the complexity coefficients estimated with each approximation method in classifying two groups of simulated time series. We also compare the computational efficiency of each approximation method. Although the basis spline method has the lowest approximation error across all simulations, we find that low approximation error does not correspond to improvements in classification accuracy. Each approximation method performs similarly in the classification task. However, the cubic spline method is the most computationally efficient and is used in subsequent chapters to estimate the complexity coefficients.

Darkhovsky and Piryatinska have shown that for a Hölder class of functions the ε -complexity coefficients capture a linear relationship between the log of the approxi-

mation error and the log of the proportion of the function used in the approximation. They have conjectured that a constant generating mechanism corresponds to constant mean complexity coefficients. In Chapter 4 we use a number of simulations of deterministic and stochastic processes to test whether, on average, the complexity coefficients are constant for a constant Hölder class of functions. For these processes, a single parameter determines both the Hölder class and the fractal dimension of the time series generated by the processes. For the tested simulations, the slope complexity coefficient B and an estimator of fractal dimension exhibit similar behavior as the parameters of the processes are varied. For three of the four processes tested, we find that the complexity coefficient B is on average, constant for a constant Hölder class of functions.

In the final chapter, the complexity coefficients are applied to the prediction of seizures in epileptic mice. Features from a 4-minute window of EEG are used to predict whether a stimulus applied to epilepsy-prone mice results in a seizure. A time series of length n can be considered as a vector in an n -dimensional space. For our dataset, 4 minutes of EEG from a single channel is represented by 2 million observations. A common method of handling such high dimensional data is to map the data to a lower dimensional set of features. For time series like EEG, features are often calculated on a uniform partition or on a sliding window taken over the length of the signal. However, EEG are inhomogeneous and this method of feature extraction may lead to features being calculated over windows with widely varying characteristics. In order to calculate our features on more homogeneous time periods, we use change points in the complexity coefficients to identify regime changes in the EEG signal. The EEG features are segmented based on change points

in the complexity coefficients and average feature values on these segments are used as predictors of a seizure response.

Models trained on features segmented in this manner are compared to several models trained on features calculated on uniform partitions of the EEG signal. The best performing models accurately classify seizure outcomes with over 80% accuracy. A model with features calculated on uniform partitions performs as well or better than the models using changes in the complexity coefficients to segment the signal. Due to the small number of trials and some inherent limitations of the data, additional tests would be needed to know how well the model performance reported here generalizes to other contexts.

Finally, we have created an R language package, `ecomplex`, that implements the ε -complexity algorithm[12]. The default settings of this package are based on the simulation experiments described in Chapter 3. Methods used for generating simulations, computing EEG features, and the classification algorithm used in Chapter 5 have been also been included in R packages to enable further study or replication of the experiments described in this work.

Chapter 2

Background

2.1 Epsilon-Complexity

The term 'complexity' has been used to describe a diverse set of both mathematical and natural phenomena. Our use of the term complexity to describe a property of a continuous is related to Kolmogorov's definition of the complexity of a discrete sequence. Kolmogorov complexity characterizes the regularity of a sequence or string by the size of the program needed to output that sequence. For example, we can take as our sequence a natural number and let our 'program' be the representation of that number in scientific notation. Then we can express 1,000,000 in 2 digits as 1E6 while the prime 7919 could not be similarly compressed [17]. For this case, 1,000,000 is less complex 7919.

The approach taken by Darkhovsky and Piryatinska in defining the ε - a complexity of a continuous function agrees with Kolmogorov's definition of complexity[5]. Roughly speaking, the complexity of a continuous function is measured by the amount of information that is needed to reconstruct that function within some error ε . The procedure used to compute ε -complexity coefficients makes successive

approximations of a function with a more sparsely sampled set of points at each iteration. The two ε -complexity coefficients are the parameters of the linear relationship between the log of the approximation error and the log of the proportion of points used for each approximation.

We begin with the formal definition of the ε -complexity of a continuous function followed by the primary theorem of [5] which shows that, for the Hölder class of functions, ε -complexity can be characterized by a pair of real numbers. To extend the definition of ε -complexity to the case of a discrete function, the function is taken to be the restriction of a continuous function to a discrete grid and the theorem is adjusted to this case. We also present the estimation procedure used to compute the ε -complexity coefficients.

Let $x(t)$ be a continuous function defined on the unit interval $\mathbb{I} = [0, 1]$. Let $\|\cdot\|$ be a norm on the function and we assume $\|x(t)\| = 1$, that is, the function has been normalized by taking $x(t)/\|x(t)\|$. Given some family of approximation methods \mathcal{F} , let $\hat{x}(\cdot)$ be an approximation of $x(t)$ reconstructed from regularly samples spaced at intervals h . Then the absolute recovery error is defined

$$\delta^{\mathcal{F}}(h) = \inf_{\hat{x} \in \mathcal{F}} \sup |x(t) - \hat{x}(t)|. \quad (2.1.1)$$

That is, $\delta^{\mathcal{F}}(h)$ is the minimum error overall all methods in \mathcal{F} for a given spacing h . We assume that $x(t)$ does not admit a perfect reconstruction. Now we define the

following as a function of the error ε :

$$h_x^*(\varepsilon, \mathcal{F}) = \begin{cases} \inf\{h \leq 1 : \delta^{\mathcal{F}}(h) > \varepsilon\} & \text{if } \{h : \delta^{\mathcal{F}}(h) > \varepsilon\} > 0 \\ 1, & \text{if no such } \hat{x}(t) \in \mathcal{F} \text{ exists.} \end{cases}$$

In words, the function $h^*(\varepsilon, \mathcal{F})$ is the minimum grid spacing h needed to approximate $x(t)$ within an absolute error not larger than ε . We can now define ε -complexity.

Definition 2.1.1. Epsilon-complexity. The number

$$\mathbb{S}_x(\varepsilon, \mathcal{F}) = -\log h_x^*(\varepsilon, \mathcal{F}) \quad (2.1.2)$$

is the $(\varepsilon, \mathcal{F})$ -complexity, or simply, ε -complexity of the function $x(\cdot)$.

We will be referring variously to the Hölder condition or Hölder continuity of a function which we now define.

Definition 2.1.2. Hölder Continuity. A function $x(t)$ is said to be Hölder continuous if there exists non-negative constants, C, α such that

$$|x(t+h) - x(t)| \leq C|h|^\alpha$$

for all t and h . We refer to this as the Hölder condition with Hölder exponent α or say a function is Hölder α .

Darkhovsky and Piryatinska have shown that for almost every individual Hölder class function, and given a rich enough family of approximation methods, the epsilon

complexity of a function has an approximately linear relationship to $\log \varepsilon$,

$$-\log h_x^*(\varepsilon, \mathcal{F}) \approx \mathbb{A} + \mathbb{B} \log \varepsilon \quad (2.1.3)$$

[5].

This relationship characterizes the ε –complexity of a Hölder function. An analogous relationship is used for the estimation of ε –complexity for discrete time series. In practical applications any function or time series is acquired as some discrete set of samples, which we will assume to be taken at regular intervals. We consider the function in the discrete case to be a continuous function restricted to a uniform grid. Then $h^*(\varepsilon, \mathcal{F})$, which we denote here $h^*(\varepsilon)$, is the proportion of sampled points, and $\frac{1}{h^*(\varepsilon)}$ is the number of sampled points, needed to approximate some function $x(t)$ within an absolute error ε . Let n be the total number of points in the initial sample, then the proportion of points needed for reconstruction with error not larger than ε is defined

$$S(\cdot) = \frac{\lfloor h^*(\varepsilon)n \rfloor}{n}.$$

Definition 2.1.3. Epsilon-complexity (discrete). The ε –complexity of a function represented by a set of values sampled on a uniform grid is $-\log S(\cdot)$.

It follows from definition 2.1.1 that

Proposition 2.1.4.

$$-\log S(\varepsilon, \mathcal{F}) \rightarrow \mathbb{S}_x(\varepsilon, \mathcal{F})$$

as $n \rightarrow \infty$.

For the discrete case, the relation analogous to 2.1.3, after substitution of $\log S$ for $\mathbb{S}_x(\varepsilon, \mathcal{F})$, is give by

$$\log \varepsilon \approx A + B \log S. \quad (2.1.4)$$

When computing the complexity coefficients, some set of grid spacings h , corresponding to the proportions S_h , are used to estimate the linear relation in 2.1.4. For each spacing h we find the minimum error approximation at that spacing. The parameters of this linear relationship, A, B , are the ε -complexity coefficients or simply the complexity coefficients.

In applications, the discrete function x_t is given by some set of samples which we assume to be uniformly sampled at an initial interval h_0 . Approximations are made using samples taken on some set of wider intervals $\{h\}$. For some set of integers $\{\alpha\}, \alpha > 1$, the function x_t is approximated using the function values downsampled to a grid with spacing $h = \alpha \cdot h_0$. These approximations are taken from one or more families of methods \mathcal{F} , for example, the family of piecewise polynomials. For a given downsample determined by $\alpha \cdot h_0$, the minimum mean-squared error of all approximation methods \mathcal{F} is then added to the set of errors, $\varepsilon_{\mathcal{F},h}$. Finally, setting $S_h = \frac{1}{h}$, a least squares linear model is fit to the set of errors $\{\varepsilon_{\mathcal{F},h}\}$ regressed on the proportion of function values $\{S_h\}$:

$$\log \varepsilon_{\mathcal{F},h} \sim A + B \log S_h.$$

The parameters of the linear model A and B are the the complexity coefficients.

Algorithm 1: Estimation of ε –complexity

Input : X a regularly sampled time series of length N .

Input : \mathcal{F} a set of approximation methods f .

Input : \mathcal{H} the set of spacings $h = \alpha \cdot h$ for $\alpha \in \{\alpha\}$.

Output: The complexity coefficients A, B .

foreach h *in* H **do**

foreach f *in* \mathcal{F} **do**

 Compute the approximation error

$\varepsilon_{h,f} \leftarrow \frac{1}{N} |f_h - X_h|^2$.

end

 Find minimum error over all f

$\text{epsilons}_h \leftarrow \min \varepsilon_{h,f}$.

end

Fit a least squares linear model

$A, B \leftarrow \text{lm}(\{S_h\} \sim \{\text{epsilons}_h\})$.

2.2 Time Series

Time series are a univariate set of discrete observations index by time. In the context of statistical modeling, time series are taken to be observations of some stochastic process. A stochastic process is a set of random variables, $X(t)$ or X_t , parametrized by time, $\{X_t : t \in \{1, 2, 3, \dots, N\}\}$. Then a time series is a sample path of a stochastic process which we denote with the lower case x_t . We note here that when considered in the context of ε –complexity, we ignore the probabilistic interpretation of a time series and instead treat the time series x_t as the restriction of a continuous function $x(t)$ to a uniform grid.

We begin by defining some standard characteristics of time series. A *strictly stationary* stochastic process is one whose distribution is independent of time t .

Definition 2.2.1. Strict stationarity. A strictly stationary time series $\{X_t, t \in$

\mathbb{Z} is a time series such that for all h

$$\{X_1, X_2, \dots, X_n\} \quad \text{and} \quad \{X_{1+h}, X_{2+h}, \dots, X_{n+h}\}$$

have the same joint distribution for all $n \geq 1$ and $h, n, h \in \mathbb{Z}$ [14].

Strict stationarity is difficult to determine in practice and we will be using the term stationary to refer to weak stationarity.

Definition 2.2.2. Weak stationarity. Where it exists, let $E(X_t) = \mu_t$ be the mean function of a time series. A time series with finite variance and mean μ_t is called *weakly stationary* if it satisfies the following conditions:

$$E(X_t^2) < \infty,$$

$$E(X_t) = \mu_t \quad \text{is a constant independent of } t,$$

$$\text{Cov}(X_t, X_{t+h}) \quad \text{is independent of } t \text{ for each } h.$$

Covariance and correlation quantify how two or more random variables jointly vary. For time series, this joint variability of a time series with itself at various time lags h is measured by the autocovariance and autocorrelation of the time series.

Definition 2.2.3. Covariance function. The *autocovariance or covariance* of a stochastic process with mean μ_t is defined

$$\sigma(h) = \text{Cov}(X_t, X_{t+h}) = E[(X_t - \mu_t)(X_{t+h} - \mu_{t+h})].$$

In the case of a zero mean function this simplifies to

$$\sigma(h) = E[(X_t)(X_{t+h})].$$

The *autocorrelation* function of the time series X_t is the normalized autocovariance and is defined

$$\rho(h) \equiv \sigma(h)/\sigma(0) \equiv \text{Corr}(X_{t+h}, X_t), \quad h \in \{1, 2, 3, \dots\}.$$

It follows that for a stationary, mean zero stochastic process with $\sigma^2 = 1$ the autocorrelation and autocovariance function are equivalent.

So far we have defined time domain properties of time series. The *spectral density* or *power spectrum* is a frequency domain decomposition of a time series. In particular, the spectral density function expresses the variance or power of a signal over a discrete or continuous frequency domain. The Wiener-Khintchine theorem relates the autocorrelation of a time series to its spectral density function.

Theorem 2.2.4. (Wiener-Khintchine theorem) For a real-valued function defined on the integers $\{t = 0, \pm 1, \pm 2, \dots\}$ then $\rho(t)$ is the autocorrelation function of a stationary time series if and only if there exists a symmetric distribution function F on $[-\pi, \pi]$

$$\rho(t) = \int_{-\pi}^{\pi} e^{it\omega} dF(\omega) d\omega \quad . \quad (2.2.1)$$

where F is called the normalized distribution function of the time series. If F has a

density function f then we call f the normalized spectral density function and

$$\rho(t) = \int_{-\pi}^{\pi} e^{it\omega} f(\omega) d\omega \quad . \quad (2.2.2)$$

[14]

To simplify, the autocorrelation function $\rho(t)$ and the spectral density function f are invertible Fourier pairs when they satisfy the necessary conditions. A time series with a slowly a decaying autocorrelation function is said to exhibit *long-range dependence*. In theory, a stochastic process with long-range dependence would not have an infinitely summable autocovariance function. In practice, a time series is given by a finite sample on which a spectral decomposition can always be performed.

Some stochastic processes can be defined by the distribution of their *increments*. The increments of a stochastic process are defined as $\{X_t - X_{t+h}\}$. The variance of the increments of a stochastic process is called the *variogram*, although the term is more commonly used in reference to spatial processes or random fields, that is, stochastic processes indexed by more than one variable.

Definition 2.2.5. Variogram. Let X_t is a zero mean stochastic process with stationary increments. The variogram of X_t is defined

$$\gamma^2(h) = \frac{1}{2} \mathbb{E} (X_t - X_{t+h})^2 .$$

Let $N(h)$ be the total of all pairs of points in a time series at distance h and

denote the size of that set by $|N(h)|$. Then the *empirical variogram* is defined

$$\hat{\gamma}(h) = \frac{1}{2|N(h)|} \sum_{i,j \in N(h)} |x_i - x_j|^2.$$

2.3 Fractal Dimension

Although fractal (or fractional) dimension is a general measure on sets, we will be considering only those sets formed by the graph of a function f or a sample path drawn from a stochastic process. Fractal dimension coincides with the usual sense of dimension for smooth manifolds; for example, the fractal dimension of a line in \mathbb{R}^1 is 1. Since we consider only graphs $\Gamma : \{(x, f(x))\}$ in \mathbb{R}^2 , the fractal dimension of these graphs should take on values between $[1, 2)$.

The Hausdorff measure of a set E , roughly speaking, measures the size of covers of that set as the diameter of covering elements $\{U_i\}$ approaches zero. We define the diameter of a set U denoted $|U|$ as $\sup\{|x - y| : x, y \in U\}$, the greatest distance between any two points in the set. A collection $\{U_i\}$ is a cover of E if $E \subset \bigcup_i^\infty U_i$ where we assume $\{U_i\}$ is countable.

Definition 2.3.1. Hausdorff Dimension. Let all elements in $\{U_i\}_\delta$ be a countable subset of $\{U_i\}$ with diameter greater than δ such that $\{U_i\}_\delta$ cover some set $E \subset \mathbb{R}^n$. The Hausdorff measure $\mathcal{H}_\delta^\alpha$ for any $\delta > 0$ is defined

$$\mathcal{H}_\delta^\alpha = \inf \left[\sum_i |U_i|^\alpha : \{U_i\} \text{ is a cover of } E \right]$$

and

$$\mathcal{H}^{\alpha(E)} = \lim_{\delta \rightarrow 0} H_{\delta}^{\alpha(E)}.$$

For $0 < \alpha < 1$, \mathcal{H}^{α} is either 0 or ∞ and there exists a critical point α at which \mathcal{H}^{α} switches from 0 to ∞ . This critical point is the Hausdorff dimension of the set E which we denote $\dim_H E$ [6].

for m do

\perp a

ny sets, the Hausdorff dimension is difficult to determine and a number of related measures are used as more tractable alternatives. The box-counting dimension is one these and, for sets in \mathbb{R}^2 , it counts the number of squares of length δ required to cover a set. It also provides an upper bound for the Hausdorff dimension of a graph. The box-counting dimension is defined as

$$\dim_B E = \lim_{\delta \rightarrow 0} \frac{\log N_{\delta}(F)}{-\log \delta}$$

where N_{δ} is the smallest number of sets that cover E of side length δ . In what follows we occasionally refer more loosely to the fractal dimension of a function or time series by which we mean, more properly, the fractal dimension of the graph of the function.

There is a general relation between the Hölder- α functions and the fractal dimension of a continuous function on \mathbb{R} . The Hölder condition is an upper bound on the box-counting dimension.

Proposition 2.3.2. Let $f : [0, 1] \rightarrow \mathbb{R}$ be a continuous function and

$$|f(t) - f(t + h)| \leq ch^{\alpha} \quad (c > 0, 0 \leq \alpha \leq 1)$$

then $\dim_B f \leq 2 - \alpha$ [6].

We use a variogram estimator of fractal dimension in Chapter 4. This method estimates fractal dimension by taking the log-log regression of the empirical variogram against increments h . This method is recommended as a robust estimator of fractal dimension by Gneiting in [9]. One justification for the estimation method is the relation between the variogram of certain stochastic processes and its fractal dimension.

Proposition 2.3.3. Let X_t be a Gaussian process with variogram satisfying

$$\gamma(h) \sim |ch|^\alpha, \quad (2.3.1)$$

then the Hausdorff dimension of the sample path of the Gaussian process, $\dim_H \Gamma(X_t)$, is almost surely $2 - \alpha$. For Gaussian processes satisfying 2.3.1, the Hölder exponent of its sample paths is almost surely α [9] [13].

2.4 Simulations

In the following two chapters we use a set of simulations to test our implementation of ε -complexity estimation and to explore properties of the ε -complexity coefficients. These simulations include both deterministic functions and stochastic processes. We introduce these processes and functions here and describe some of their properties.

The *logistic map* is a single parameter nonlinear deterministic discrete difference equation defined

$$x_t = r(x_{t-1})(1 - x_{t-1}).$$

For values of the parameter r above approximately 3.6 but less than 4, the trajectory of the logistic map exhibits chaotic behavior and the graph of the function changes significantly for small perturbations of the r parameter.

The *ARMA* model is a widely used stochastic model that combines autoregressive (AR) and moving average (MA) models. The AR component expresses an observation at time t as a linear combination of preceding observations plus a random error component:

$$X_t = b_1 X_{t-1} + \cdots + b_p X_{t-p} + \varepsilon_t$$

The number of previous observations p is the order of the model, denoted $AR(p)$. The random error, ε_t , is assumed to be independent, identically distributed (IID) observations from a standard normal distribution with mean 0, variance σ^2 and covariance $Cov(X_i, X_j) = 0, i \neq j$. The latter is called white noise or Gaussian noise and the errors are said to be drawn from a white noise process $\{\varepsilon_{t_i}, \varepsilon_{t_j}\} \sim WN(0, \sigma^2)$.

The moving average(MA) process of order q models the a time series as a linear combination of q samples $\{\varepsilon_t\} \sim WN(0, \sigma^2)$

$$X_t = \varepsilon_t + a_1 \varepsilon_{t-1} + \cdots + a_q \varepsilon_{t-q}. \quad (2.4.1)$$

For our simulations we use parameters that define stationary $ARMA(p, q)$ processes.

A stationary ARMA process has an autocorrelation function $\rho(h)$ that decays at an exponential rate as $|h| \rightarrow \infty$ [14]. We include several functions that have a slowly decaying autocorrelation function, a characteristic of processes with long-range dependence. The first of these is the fractionally integrated ARMA process or

FARIMA process. The ARIMA process is model whose integral differences d are a ARMA(p,q) processes, denoted ARIMA(p, d, q). The FARIMA model modifies the ARIMA process by allowing for fractional values d [14].

The *Weierstrass* function is a well-known example of a continuous but nowhere differentiable function. The Weierstrass function can be written in several forms. Here we use the form that isolates the parameter α corresponding to the Hölder exponent of the function.

Definition 2.4.1. Weierstrass function. The *Weierstrass* function defined

$$W_\alpha(x) = \sum_{n=0}^{\infty} b^{-n\alpha} \cos(b^n \pi x) \quad (2.4.2)$$

is a continuous periodic nowhere differentiable function that is Hölder- α :

$$|W_\alpha(x) - W_\alpha(y)| \leq C |x - y|^\alpha.$$

The parameter α determines both the fractal and Hausdorff dimension of the Weierstrass function. The Weierstrass function has box-counting dimension equal to $2 - \alpha$, providing the upper bound for the Hausdorff dimension [6]. That the lower bound of the Hausdorff dimension is also equal to $2 - \alpha$ was only recently proved by Shen [15].

A variant of the Weierstrass function is *the random-phase Weierstrass* function. The random phase function has the same form as the Weierstrass function with an added phase, ϕ , drawn from a uniform distribution:

$$W_\alpha(x) = b^{-n\alpha} \cos(b^n \pi x + \phi_n), \quad \phi_n \sim \text{unif}(0, 1).$$

The Hausdorff dimension of the random phase variant Weierstrass function has also been proved to be $2 - \alpha$ [10]. See Figure 4.3 in the following chapter for illustrations of both Weierstrass functions.

Brownian motion and *fractional Brownian motion* (fBm) are both Gaussian processes that can be defined in terms of their increment process. Brownian motion has normally, independently distributed increments and fractal dimension $1\frac{1}{2}$. On the other hand, fBm has *dependent* increments with

$$E[(X_t - X_{t+h})^2] = h^{2\alpha}.$$

When $\alpha = \frac{1}{2}$ fBm is equivalent to Brownian motion.

For fBM, the parameter α determines both the fractal dimension, Hurst parameter, and the Hölder exponent of the sample paths. When $\alpha > \frac{1}{2}$ fBm has lower fractal dimension and exhibits long-range dependence with α corresponding to the Hurst coefficient of the process. When $\alpha < \frac{1}{2}$ sample paths of fBM have a higher fractal dimension and a lower Hurst parameter.

In contrast to fBm, the *Cauchy Process*, has two parameters which separately control the long-range dependence of the function and its local behavior or fractal dimension [8]. The Cauchy process is a stationary Gaussian process determined by its autocorrelation function:

$$\rho(t) = (1 - |ct|^\alpha)^{-\beta/\alpha}, \quad \alpha \in (0, 2]; \beta > 0.$$

The Cauchy process parameter α determines the fractal dimension D of the process

via

$$D = n + 1 - \frac{\alpha}{2},$$

and for $\beta \in (0, 1)$ the Cauchy process has long-range dependence with Hurst parameter

$$H = 1 - \frac{\beta}{2}.$$

2.5 Approximation Methods

In calculating the complexity coefficients we use some set of approximation methods \mathcal{F} . In the initial implementation, the complexity coefficients were estimated using piecewise polynomials of up to degree 5. For our implementation, we added several methods to the procedure for estimating ε -complexity. In Chapter 4, we compare the accuracy of these approximation methods on our set of simulations. The ε -complexity coefficients are estimated for two groups of simulations using each approximation method. We compare both the average minimum approximation errors of each method and the classification accuracy of classifiers trained on the complexity coefficients as estimated by each method. Here we briefly describe the three approximation methods and some of the properties.

Splines are piecewise polynomials joined on some set of knots, or intersection points, with the additional constraint that the spline function of order n has continuous derivatives of order $n - 1$ at each knot. Two spline methods were used in our implementation, cubic splines and basis splines of order up to 5.

Basis or B-splines of order n are smoothly joined splines that form a basis for the space of n -order splines. That is, any n -order spline can be written as a unique linear

combination of B-splines. For regression splines, an intermediate design matrix is used in the calculation of the final approximation. This means the computational complexity and memory requirements of the method can grow at a greater than the quadratic rate with the number of basis elements used.

The cubic spline method is a spline method but the splines are restricted to degree 3 piecewise polynomials. For cubic splines, the design matrix is tridiagonal, having entries only on its diagonal and two off-diagonals. An LU type decomposition for tridiagonal matrices allows for linear time cubic spline interpolation [4]. In practice, this means cubic spline interpolation with a fixed number of knots is much less costly in terms of both memory and computational efficiency compared to B-spline regression methods that use matrix storage and matrix arithmetic.

For a third method we implemented an interpolating subdivision technique which we will refer to as the lifting scheme after Sweldens. Sweldens gives a practical introduction to interpolating subdivision in [16]. To convey the basic approach, we consider a simple example of repeated interpolation. Given a set of regularly sampled points $\{s_{0,k}\}$ with $k \in \mathbb{Z}$ points a linear interpolation of the point $x_{j+1,2k+1}$ is the average of the adjacent points $s_{j,k}$ and $s_{j,k+1}$. A recursive formula for averaging and interpolating mid-points at stage $j + 1$ is given by

$$\begin{aligned} s_{j+1,2k} &= s_{j,k} \\ s_{j+1,2k+1} &= 1/2(s_{j,k} + s_{j,k+1}). \end{aligned}$$

For this example, repeated iterations would simply interpolate two points with a line. If the sampled function was linear, the interpolation would be an exact recon-

struction.

The lifting scheme is similar to wavelet interpolation methods. Unlike wavelet based interpolation, though, the lifting scheme can be extended to interpolating irregularly spaced points since uniform spacing is not a condition for polynomial interpolation. For the ε -complexity algorithm the pattern of interpolation changes based on the grid spacing h at which samples are taken. However, for a set interpolation pattern the polynomial interpolant needs to be computed only once. Neville's algorithm is used to compute a set of $2 * (n - 1)$ weights for an n -order polynomial interpolation. To use the simple example order 2 or linear interpolation above, the weights to compute $x_{j+1,2k}$ at the mid-point of $s_{j,k}, s_{j,k+1}$ are $(1/2, 1/2)$. Then $x_{j+1,2k}$ is calculated

$$\langle 1/2, 1/2 \rangle \cdot \langle s_{j,k}, s_{j,k+1} \rangle = x_{j+1,2k}.$$

These weights, $(1/2, 1/2)$, can be thought of as a time-domain filter. For our approximation, similar filters were computed for each order polynomial and interpolation pattern. Interpolation of a given point is then accomplished with a similar inner product, with adjustments made for boundary points.

For our implementation polynomials of orders 1, 3, 5 were used for interpolation and for each order polynomial three filters were computed for each distinct interpolation pattern. Since these filters needed only to be generated once, the interpolation method is $O(n)$ or linear in the number of inputs.

Chapter 3

Approximation Methods

In theory, the ε –complexity of a continuous function is found using the minimum error over a possibly infinite set of approximation methods. To be computationally efficient, this set of approximation methods needs to be finite and in practice, is restricted to some relatively small set of approximation methods. The initial implementation of the ε –complexity algorithm used piece-wise polynomial interpolation. Our implementation extends the number of methods used to three: basis splines(B-splines) of up to order 5, cubic splines and interpolating subdivision. A fourth method takes the minimum over all methods at each stage of the estimation procedure.

By drawing on a richer set of reconstruction procedures we hoped to improve approximation accuracy and the performance of the ε –complexity coefficients. In this chapter, we gauge the approximation methods in these two ways. First, we measure the average approximation error of each method on a range of functions and stochastic processes. We then test the ability of the complexity coefficients to discriminate between two sets of simulations with similar parameters. Based on these experiments, increasing the accuracy of the approximation did not lead

to better performance in the discrimination task. In fact, the cubic spline method which was less accurate in terms of approximation errors, performed as well or better than the more accurate B-spline method.

A separate concern was the computational efficiency of the approximation methods. Computational efficiency becomes a pressing issue as the size of a data set increases. Although the lifting scheme and cubic spline approximation are both computationally linear in their inputs, the cubic spline implementation was the fastest method by an order of magnitude. Based on computational efficiency and classification performance, we have used the cubic spline method for computing the complexity coefficients in our R implementation and in the applied work in the following chapters.

3.1 Approximation and Classification Accuracy

For these experiments, we used six simulations methods comprised of both deterministic and stochastic processes. Two groups of processes were used each with slightly varying initial parameters. The processes and the starting parameter values for each group are listed in Table 3.1. For each process, 30 samples were generated by drawing a parameter uniformly from a small window centered on the initial parameter values. A single sample observation from each simulation group is shown in Figure 3.1.

For each of the approximation methods we used a simple diagnostic plot to check that the expected log-log relation of approximation errors ε_h to the proportion of points S_h used for reconstruction was roughly linear. Shown in 3.2 is an example

Process	Parameters	Group 1	Group 2
ARMA	ϕ, θ	(-0.1, 0.3), (0.2, 0.1)	(0.4, 0.3), (0.1, -0.5)
Cauchy	α, β	0.1, 0.5	0.3, 0.7
FARIMA	ϕ, d, θ	(0.1, -0.5), 0.3, (0.6, 0.01)	(0.2, -0.4), 0.3, (0.4, 0.02)
fBm	α	0.1	0.3
Logistic	r	3.87	3.70
Rand. Weierstrass	α	0.65	0.17

Table 3.1: Initial parameters for each simulation group.

of log-log linear fit of the approximation errors against the sample proportion for a single function from each simulation type. This figure shows the regression plots generated by the lifting method and the B-spline and cubic spline methods resulted in similar, approximately linear, regression plots.

A plot of the complexity coefficients generated by each approximation method and group is shown in Figure 3.3. Each point represents where an individual function lies in the feature space of the two complexity coefficients A and B . The coefficient values have been normalized to the $[0, 1]$ interval so the plots show the relative distribution of the functions in the feature space. Although we will quantify the ability of the approximation methods to separate the two simulation groups below, the scatter plot shows that the approximation methods place each process at similar coordinates in the complexity coefficient space. The scatter plots also show that some methods separated particular processes better. For example, the cubic spline method separates the two groups generated by ARMA(2,2) processes, as seen on the middle left-hand side of Figure 3.3.

The scatter plot of the simulations also provides some information about the direction in which the simulations are separable. If the simulations from each group

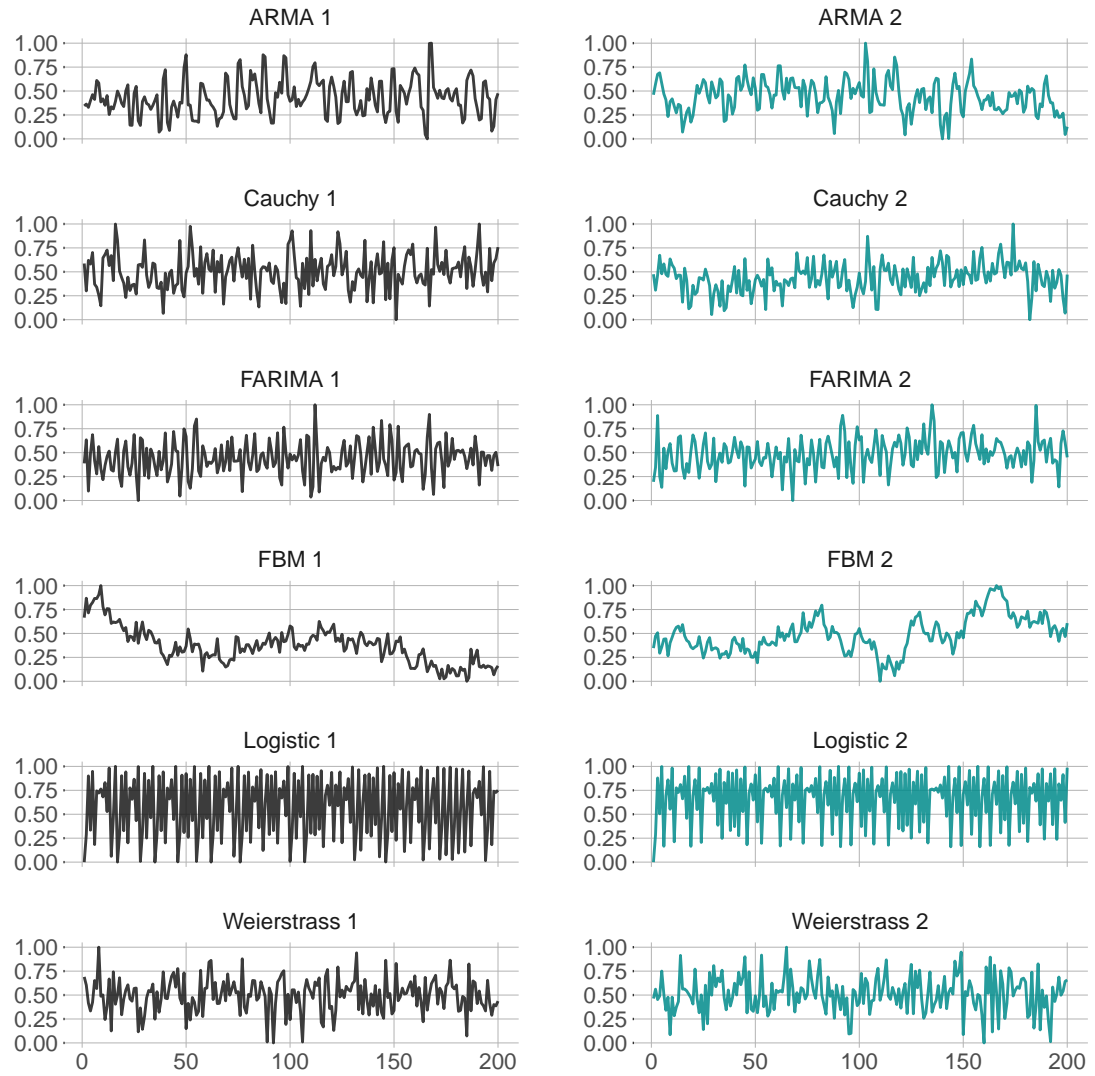


Figure 3.1: Sample functions from each simulation group.

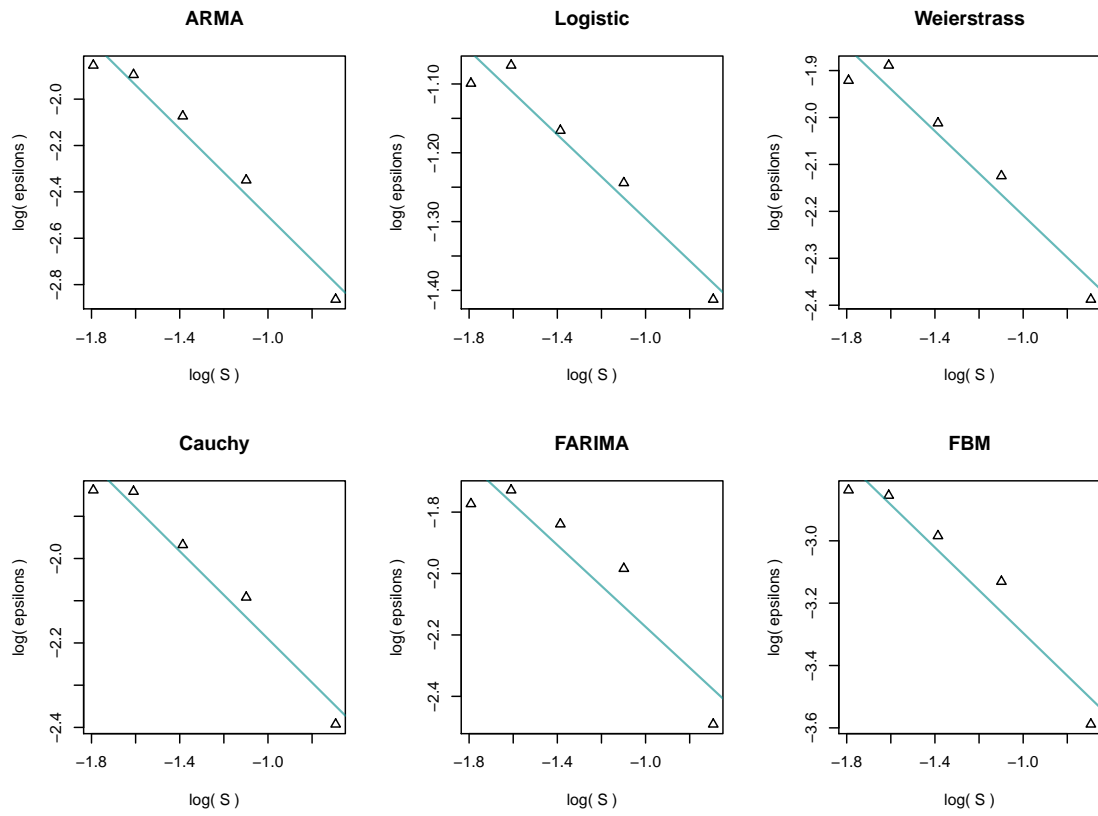


Figure 3.2: Linear regression plots for the lifting scheme.




Figure 3.3: The two simulation groups plotted in the complexity coefficient space.

were separable along a single axis, either vertically or horizontally, this would indicate that a single complexity coefficient could be used to separate the groups. Figure 3.3 shows that for the processes that are readily separable —fBm and the ARMA and logistic simulations —both the slope and intercept coefficient appear to be useful in separating the two sets of simulations. For example, both the fBm and ARMA simulations are well separated along the x -axis, corresponding to the complexity coefficient A . On the other hand, the realizations of the logistic map in the upper right hand corner appears separable on the diagonal meaning both complexity coefficients contributed to differentiating the two simulation groups.

Using the first set of simulations described above, 30 samples were generated and a total approximation error for each method was calculated. The approximation error was calculated by taking a simple sum of errors at each down sampling level

h :

$$\varepsilon_{\mathcal{F}} = \sum_h \varepsilon_{h,\mathcal{F}}.$$

Table 3.2 shows the mean approximation error over the 30 samples. The B-spline method had the lowest MSE for all processes. Since the B-spline method produced the minimum error for each function, the combined method simply reflects that of the B-spline approximation. The lifting method produced a slightly lower mean error but the cubic spline and lifting errors are very similar for each simulation.

Function	Lift	Cspline	Bspline	Combined
ARMA	0.10	0.11	0.07	0.07
Cauchy	0.09	0.10	0.06	0.06
FARIMA	0.13	0.14	0.08	0.08
fBm	0.03	0.04	0.02	0.02
Logistic	0.30	0.32	0.20	0.20
Weierstrass	0.14	0.15	0.09	0.09

Table 3.2: Mean total approximation error for 30 samples of each simulation.

Function	Lift	Cspline	Bpspline	Combined
ARMA	0.03	0.00	0.07	0.07
Cauchy	0.65	0.53	0.52	0.53
FARIMA	0.32	0.30	0.32	0.33
fBm	0.20	0.22	0.20	0.22
Logistic	0.00	0.00	0.00	0.02
Weierstrass	0.50	0.43	0.45	0.45

Table 3.3: Classification errors using complexity coefficients A and B for each approximation method.

The set of complexity coefficients as computed by each approximation method were then used to classify the 30 samples from each simulations. The classifier was

a random forest with 500 trees with 3 features used to determine branches. The overall out-of-bag classification error is reported in Table 3.3. The cubic spline method performed equally well or better on 4 of the 6 methods, but the results of all methods are similar.

3.2 Computational Efficiency

For our set of processes, classification performance did not appear adversely affected by the slightly lower approximation accuracy of the lifting and cubic spline methods. However, both the cubic spline and lifting approximation methods are computationally more efficient than the B-spline method. Both methods are linear in the number of inputs and both are linear or constant in terms of their space complexity, that is, the amount of memory needed to store intermediate computations is linear or constant function of the size of the inputs.

The execution time of each approximation method was estimated using the system time taken to complete approximations on inputs of size 10^7 to 10^{13} . The average computational time for each input size on a log scale is shown in Figure 3.4. For our implementation, the B-spline method increases the number of knots linearly with the inputs and the associated design matrix grows quadratically with the number of inputs. This leads to the B-spline method becoming infeasible for even moderately large inputs. Although the cubic spline and lifting method are computationally linear, the lifting scheme was implemented entirely in the R language. The cubic spline method is called from R but most calculations are calculated in the more computationally efficient C language. Although it is not clear in Figure 3.4,

the cubic spline method was an order of magnitude faster than the lifting scheme.




Figure 3.4: Computational time of each approximatio method as a function of the size of the log of the input.

Our primary goal of the preceding test was to determine if enlarging the set of approximations would lead to improved performance of the complexity coefficients. We assumed that the enlarged set would allow for improved approximation error and a better estimation of the complexity coefficients. Here we measured performance by the ability of the complexity coefficients to discriminate between closely related simulations. Lower approximation errors did not correspond to better classification accuracy, however. The cubic spline method was the most computationally efficient and performed as well or better on the classification test as the other methods.

Based on these results, the cubic spline approximation was used for the applied work in Chapters 4 and 5.

Chapter 4

Interpreting the Complexity Coefficients

In this chapter, we explore how the complexity coefficients relate to other characteristics of time series. In particular, we use a number of simulations to examine how the complexity coefficients change as both the fixed Hölder- α class of a simulation and its fractal dimension is varied. For these simulations a single parameter controls both the Hölder class and the fractal dimension of the simulations. We find that the slope complexity coefficient B and a fractal estimator change linearly with this parameter for all but one of the simulations.

We begin by looking at the behavior of the complexity coefficients for a few simple functions with added noise. The complexity coefficients are given by parameters of $\log - \log$ linear fit to points whose original values is less than 1 and the interpretation of the parameters may not be intuitive. These examples highlight some basic properties of the complexity coefficients.

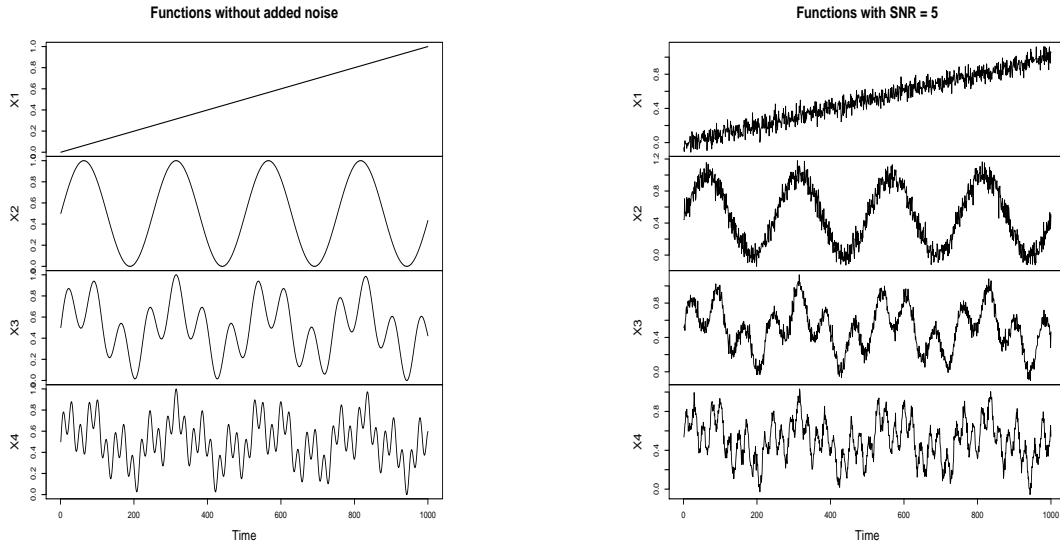


Figure 4.1: Simple linear and sinusoidal functions (left) and with added noise (right).

4.1 The Complexity Coefficients

The epsilon complexity coefficients A and B along with a variogram-based fractal dimension estimator, which we denote \hat{D} , were estimated on the four functions depicted in Figure 4.1. In addition to the original functions, A, B and \hat{D} were computed on the functions with two levels of added noise. Figure 4.2 shows the *log-log* regression of the errors ε_h on the fraction of samples kept S_h . The complexity coefficients for these fits are given in Table 4.1. While these functions do not reveal much about the behavior of the complexity coefficients for more complicated time series, the results demonstrate some simple properties of the complexity coefficients and the sensitivity of the estimators to added noise.

The parameters of the linear regression in Figure 4.2 determine the complexity coefficients A, B . The x -axis is the log of S_h , the proportion of points used to approximate the original function at each step. The finest scale approximation is




Figure 4.2: The log-log linear regression of approximation errors against sample proportion S_h for simple linear and sinusoidal functions.

Without Noise	A	B	\hat{D}
x	-41.45	-1.92	1.00
$\sin 5x$	-22.71	-4.33	1.00
$\sin 5x + \sin 17x$	-18.71	-4.61	1.00
$\sin 5x + \sin 17x + \sin 53x$	-14.54	-4.81	1.01
SNR = 20	A	B	\hat{D}
x	-4.04	-0.57	1.97
$\sin 5x$	-3.90	-0.59	1.94
$\sin 5x + \sin 17x$	-4.07	-0.55	1.94
$\sin 5x + \sin 17x + \sin 53x$	-4.12	-0.52	1.67
SNR = 5	A	B	\hat{D}
x	-3.39	-0.54	1.99
$\sin 5x$	-3.48	-0.55	1.99
$\sin 5x + \sin 17x$	-3.58	-0.55	1.99
$\sin 5x + \sin 17x + \sin 53x$	-3.70	-0.58	1.82

Table 4.1: Complexity coefficients and fractal dimension estimates for linear and sinusoidal functions.

when $h = 2$, represented by the point with the least (log) error located furthest to the right at $\log(S) \approx -0.7$. Small values at this point correspond to accurate

approximations and less variability at the finest scale of the sampled function. In theory, the approximation $\log(1) = \text{zero}$ should be an exact approximation of the function. Since we are computing the log of the approximation error ε , $\log(\varepsilon) \rightarrow -\infty$ as $\varepsilon \rightarrow 0$ and an intercept of -41.45 represent the log of an error that is approximately zero, give or take some machine accuracy. For the three functions in Figure 4.1 the overall approximation error is low. Larger intercept should correspond to larger errors and, as expected, the intercept value A increases with increased function complexity and higher levels of noise.

The complexity coefficient B measures the rate of change of the approximation error as a function is approximated using fewer samples. The coefficient B increases rapidly as noise is added. Similarly, the fractal dimension estimator approaches its theoretical maximum 2 as Gaussian noise is added to the functions. In order to compare these results to estimates on pure Gaussian noise, we estimated the complexity coefficients and fractal dimension on 50 samples Gaussian noise. The mean estimate of the complexity coefficient B was -0.54 , close to the estimates in Table 4.1 for functions with added noise. For both the fractal dimension estimator and the complexity coefficient B , relatively small amounts of added noise generate estimates similar to that of Gaussian noise.

	A	B	\hat{D}
White Noise	-2.84	-0.54	2

Table 4.2: Complexity coefficients and fractal dimension estimates for Gaussian noise.

4.2 Hölder Class and Fractal Dimension

Several of the simulations used in Chapter 3 to test the performance of approximation methods have parameters that determine both the Hölder class of a function and its fractal dimension. Our experiments show that for most of these simulations, the complexity slope coefficient B has behavior like the variogram estimator of fractal dimension \hat{D} . For three of the simulations, the Weierstrass function, fractional Brownian motion (fBm), and the Cauchy process, the median values of the B and \hat{D} change linearly with the parameters determining both fractal dimension and the Hölder exponent of the function or sample paths of the simulations. For these simulations, the complexity intercept coefficient A changed non-linearly as the parameter determining the Hölder class of a function changed. A different pattern was observed for the random-phase Weierstrass function. For samples of the random-phase Weierstrass function, the values for the complexity coefficient B and \hat{D} were similar to white noise while the complexity coefficient A changed linearly with the parameter determining the Hölder exponent of the function.

The Weierstrass function is a self-similar deterministic function whose Hölder exponent and fractal dimension when written

$$W_\alpha(x) = \sum_{n=0}^{\infty} b^{-n\alpha} \cos(b^n \pi x) \quad (4.2.1)$$

is determined by the single parameter α . The Hölder exponent is equal to α while the function's fractal dimension is $D = 2 - \alpha$. Since the function is deterministic, the complexity coefficients and fractal estimator were calculated on a single example of length 1000 for 20 parameters of α in the interval $(0, 1)$. The results in Figure

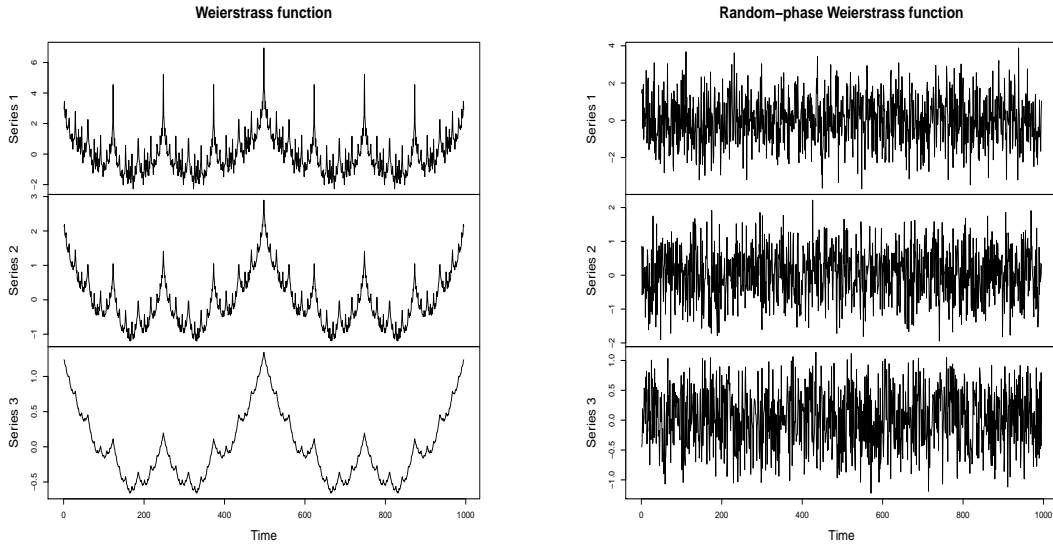


Figure 4.3: The Weierstrass and random-phase Weierstrass function for $\alpha = \{0.20, 0.42, 0.80\}$

4.4 show that the complexity coefficient B and fractal dimension change linearly with α . The fractal dimension estimator tracks the theoretical fractal dimension closely, ranging from 2 to 1. The complexity coefficient B ranges -0.6 to -1.4 . For both estimators, functions with a theoretical fractal dimension near 2 produced estimates similar to those produced by fractal and B complexity coefficient estimates for Gaussian noise. This suggests a value of around -0.5 is near an upper bound on the level of complexity or randomness that the B coefficient can measure.

The random-phase Weierstrass function adds a random phase to the periodic component of the Weierstrass function but its Hölder exponent and fractal dimension are determined by the α parameter in the same manner as for the deterministic Weierstrass function. The Weierstrass and random-phase Weierstrass functions are illustrated in Figure 4.3 for three different parameter values $\alpha = \{0.20, 0.42, 0.80\}$. Unlike the other results reported in this Chapter, neither the fractal estimator nor




Figure 4.4: Complexity coefficients and fractal dimension for values of the Weierstrass α parameter.

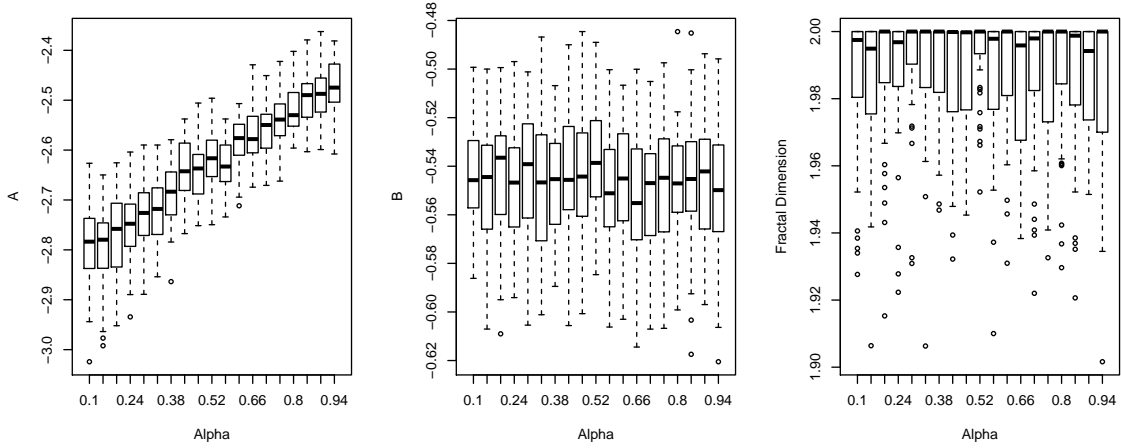


Figure 4.5: Complexity coefficients and fractal dimension for values of the random-phase Weierstrass α parameters.

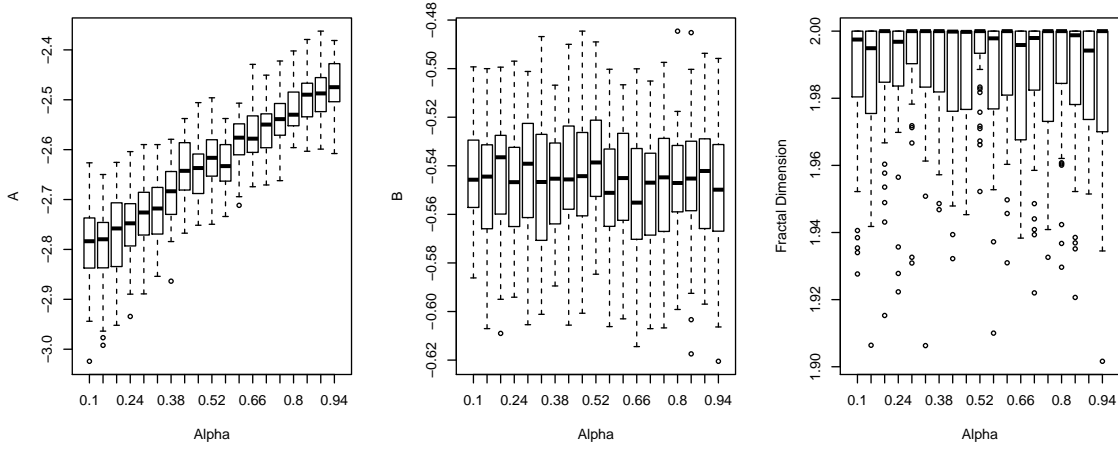


Figure 4.6: Complexity coefficients and fractal dimension for values of the random-phase Weierstrass α parameters.

the complexity coefficient B varied with α . For each parameter the complexity coefficients and fractal dimension were estimated on 50 samples of length 1000. Figure 4.6 shows that the complexity coefficient B and fractal estimator \hat{D} range around values associated with estimates for Gaussian noise. The intercept coefficient A increases linearly but the total magnitude of the change is relatively small—the median of the estimates has a range of less than 0.5.

Fractional Brownian motion also has a single parameter α which is equal to the Hölder exponent and determines the fractal dimension of sample paths as $D = 2 - \alpha$. For fBM the parameter α , sometimes denoted H , is also the Hurst parameter. The results of the estimators computed on 50 samples are shown in Figure 4.8 and are similar to the estimates for the deterministic Weierstrass function with the median of both the complexity coefficient B and \hat{D} changing linearly with the parameter α . As was the case for the deterministic Weierstrass function, the complexity coefficient

A also changes non-linearly with α .

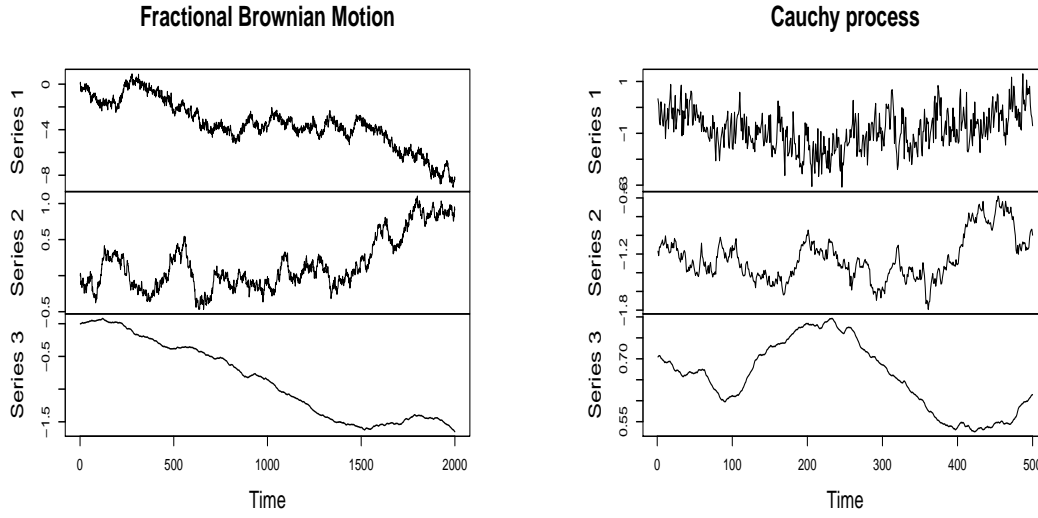


Figure 4.7: Fractional Brownian motion and the Cauchy process with $\alpha = \{0.20, 0.42, 0.80\}$ for a constant β parameter.

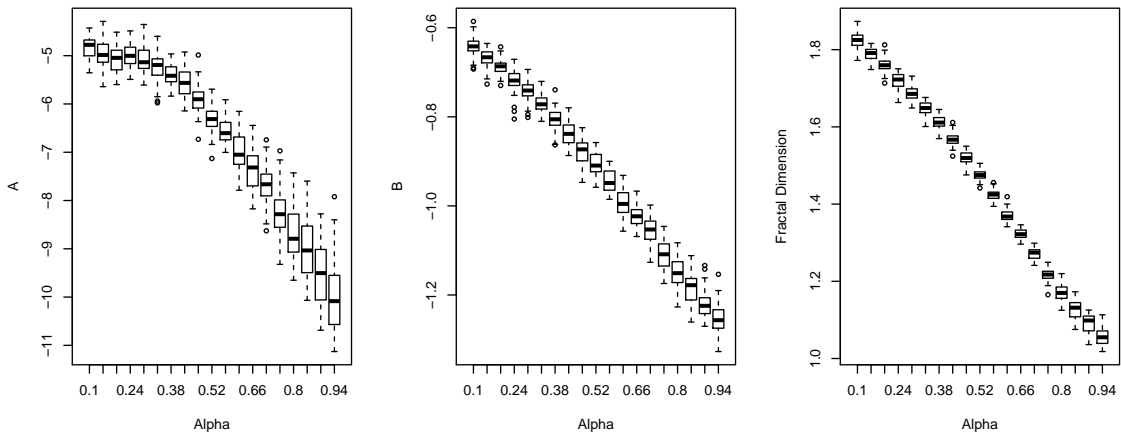


Figure 4.8: Complexity coefficients and fractal dimension for values of the α parameter of fractional Brownian motion.




Figure 4.9: Complexity coefficients and fractal dimension for various α parameters for a constant β parameter of the Cauchy process.

The Cauchy process was the final simulation we tested. The Cauchy process has two parameters: the parameter α determines fractal dimension and the parameter β determines the Hurst coefficient or long-range dependence of the sample paths. Simulations were generated on a grid of the parameters α, β for $0 < \alpha < 2$ and $0 < \beta < 1.5$ with 30 simulations generated for each point on the parameter grid. At a constant parameter β the median of the complexity coefficients and fractal estimator again show a linear relation to the α parameter, as shown in Figure 4.9. This relation holds even as β , the Hurst parameter, is varied. The mean value of each of the estimators for a fixed value β and α is shown in Figure 4.10. Variation in the Hurst parameter has little affect on either the complexity coefficients or the fractal dimension estimator.

Figure 4.11 shows another view of the complexity coefficients plotted against changes in the Cauchy process parameter β . The complexity coefficient B and fractal




Figure 4.10: Mean value of the complexity coefficients and fractal dimension for each α value of the Cauchy process.

dimension \hat{D} , preserve the grid of the parameter space—for a fixed β , the parameter varies linearly with α . The β parameter controls the long-range dependence of the Cauchy process and the plots indicate that the complexity coefficients, like fractal dimension, measure a local, property of a function.

For all simulations but the random-phase Weierstrass function the, a similar pattern was observed for the complexity coefficients and the fractal dimension estimator. Both the fractal dimension estimator \hat{D} and the complexity coefficient B changed linearly with the parameter determining the Hölder exponent and fractal dimension of the simulations. On the other hand, our initial example of simple functions and the random-phase Weierstrass functions appears shows a relationship between the intercept coefficient A and fine-scale noise. These results are for a small set of processes whose fractal dimension and Hölder exponent are determined by a single paramater. A comparison of the complexity coefficient B and fractal




Figure 4.11: The mean value of the complexity coefficients and fractal dimension for each β value of the Cauchy process.

estimators on a wider variety of time series might show whether the relation the ε -complexity coefficients and fractal dimension hold for a wider range conditions.

Chapter 5

Seizure Prediction

Electroencephalograms (EEGs) capture the electromagnetic potential of aligned neurons firing in unison. Although a number of correlations between EEG and physiological phenomena have been found the relation between EEG and the dynamics of local neural circuitry are not well-understood [3]. One difficulty is relating in vitro behavior of individual neural cells or small clusters of cells with the more complex large scale variations recorded by the EEG signal. In the case of the epilepsy, the challenge is to relate better known cellular mechanisms to the spread of synchronized neuronal firing associated with epileptic seizures. While there have been numerous studies on the prediction of naturally occurring seizure predictions, this study attempts to predict what brain states correlate with the ability to *induce* a seizure. This allows us to compare conditions that are conducive to or inhibit a seizure reaction.

In this chapter we present an exploratory study to determine whether a seizure response to a locally applied stimulus can be predicted based EEG recorded prior to the stimulus. Prediction of seizure outcomes is made using EEG taken from both local field potential (LFP) electrodes located in the thalamus that capture the firing

of a small number of neurons and intercranial EEG electrodes located at the surface of the brain that capture the aggregate behavior of larger clusters of neurons. While we do not track the variation of the EEG signal in time, predictions made on each EEG channel allow us to compare differences in EEG across brain regions.

Using a 4 minute window of EEG taken from an epilepsy prone mouse we predict whether a stimulus will induce a seizure. A small set of spectral and non-linear features are calculated on six EEG channels. Change points in the ε -complexity coefficients are used to segment the features and the results are compared to other models that partition the features uniformly. Each of six model is trained on each of the six EEG channels, resulting in 36 models in total. This approach allows us to compare how features predictive of seizure outcomes vary by brain region and how this variation changes based on the segmentation model used. In particular, we find an increase in gamma and theta band power in the LFP channels is predictive of seizure outcomes across all models. We also find a distinct pattern of the variation in feature importance across channels. This gives us some insight into how EEG characteristics predictive of seizures varies by channel.

5.1 Introduction

EEG are non-stationary signals marked by transient waveforms and regime changes over relatively short periods of times. A common method of analyzing EEG is to compute some set of features on the raw signal EEG signal and use these features as the input to one or more classifiers. These features often combine signal decomposition, for example, spectral, wavelet, principal component(PCA) or independent

component(ICA) decompositions, with some linear or non-linear features[1]. Time series like EEG are high dimensional vectors: signal decomposition and feature selection serve to reduce the dimension of the EEG signal and result in a more interpretable feature set. Dimension reduction also entails simplifying the signal by, for example, averaging features over a sliding window. A drawback to this method is that averaging over arbitrary windows may obscure important features associated with the varying dynamics of the brain.

In our approach, we attempt to capture intrinsic variations in the signal by computing features on more homogeneous segments of the EEG signal. We use changes in the ε -complexity coefficients to detect discontinuities or regime changes in the EEG signal and these change points determine how features are segmented. We compare the performance of a classifier using features segmented on complexity coefficient change points, to classifiers trained on uniformly partitioned features.

5.2 EEG Data

The EEG data were gathered from 4 mice with a genetic mutation in the voltage-gated sodium channel gene *Scn1a*. A similar mutation is responsible for the Dravet syndrome and the mutation results in early onset epileptic seizures [11]. The mice were equipped with 4 intracranial(iEEG) electrodes —intracranial electrodes placed directly on the brain cortex —along with 2 local field potential (LFP) electrodes located in the thalamus. The mice were genetically modified to produce a light sensitive opsin protein. This allows for direct stimulation of neurons with a laser pulse train designed to induce a seizure in the mouse.

Data was gathered from 13 distinct time periods during which the stimulus was applied between one to four times. The 4-minute segments preceding the stimulus was used to predict outcomes and there were 26 trials in total. The result of each trial was were coded as either a seizure or no seizure and we refer to a seizure as a positive response. Coding was based on the visible response of the mouse. The EEG signals around the stimulus period were also visually examined to check the accuracy of the coding. The iEEG and LFP signals was sampled at 1220.7 Hz and a bandpass filter removed frequencies below 0.5 Hz and greater than 300 Hz. A notch filter was applied at 60 Hz and its harmonic frequencies.

5.3 EEG Features

The features used to predict seizure outcomes consisted of windowed spectral band power estimates, variance, spectral entropy and the Hurst parameter. All features were calculated on non-overlapping 2 second intervals. The power in frequency bands delta, theta, alpha, beta, and gamma was calculated as the integral of the spectral density $f(\lambda)$ computed as a smoothed periodogram. For example, delta band power, δB , corresponds to the frequency band $0.5 - 4Hz$ and integrates $f(\lambda)$ over over this interval

$$\delta B = \int_{0.5}^4 f(\lambda) d\lambda.$$

The frequency bins corresponding to the remaining theta, alpha, beta and gamma frequency bands are $4 - 8Hz$, $8 - 12Hz$, $12 - 30Hz$, $30 - 100Hz$, respectively. Band power for each interval was normalized by dividing by total band power on the interval

Spectral entropy is a measure of the distribution of power in the spectrum of a signal.

$$SE = - \int_{-\pi}^{\pi} f(\lambda) \log f(\lambda) d\lambda.$$

This calculation is similar to the calculation of information entropy where a frequency bin takes place of a discrete event. We have defined the *Hurst parameter* in Chapter 2. The Hurst parameter is a measure of long-range of a time series and indicates a slow decay in a time series autocorrelation function. The Hurst parameter was estimated using the corrected empirical Hurst exponent[18].

5.4 Classification Procedure

Prediction of a seizure outcome is treated as a binary classification task with the two classes being a seizure response or no seizure response. The classification procedure takes place in three main steps: The time series is segmented based on one of our six partition methods, the classifier is trained on these segments, and new observations are assigned a class probability based on the weighted sum of probabilities assigned by the trained classifier. Classifier performance was evaluated using 5-fold cross-validation and positive responses were evenly spread across the 5 test sets. The full classification procedure for segmentation based on ε -complexity is outlined below.

1. For each channel compute set of features including ε -complexity on regular intervals.
2. For each channel, compute the change points in the ε -complexity coefficients.
3. For each trial, segment all features using the change point set computed in the

previous step.

4. Compute the mean of each feature on the segments.
5. Label the means of the segments in a training set with the label of the full time series and use this set to train a classifier.
6. Compute the mean value of each feature for each segment.
7. Train a classifier on the segmente feature meeanas.
8. Using the held-out test set, compute the class probability of each segment of the test trial using the trained classifier.
9. For each trial the sum of the prediction probabilities of each segment, weighted by segment length, determine the final class probability.

In the case of a uniform partition of the features, this method is equivalent to simply using the average class probability for each segment to classify the trial.

In theory, any single classifier or several classifiers could be used in steps 5 and 6. For study we use a single classifier, random forest, which is known to perform well in a wide range of contexts. It has also been used in earlier studies employing the complexity coefficients as features. For pure prediction performance an ensemble of classifiers would likely perform better but the results might less interpretable.

The random forest classifier is constructed from a large number of individual decision trees. For numeric data with d features, the decision trees partition the d -dimensional feature space into d -dimensional rectangles or hyperrectangles, each labeled with a class[7]. Each decision tree's branches correspond to partitions of the

feature space. For each branch of a decision tree, a subset of the features are chosen, restricting the decision to some subspace of the feature space, and the division of that subspace is made in order to increase the class purity of the partition. The final prediction is based on the combined vote of the individual trees. In addition to a random selection of features, a random subset of observations are used to build each tree. This allows for an out-of-bag(OOB) estimate of the classifier accuracy which is calculated by classifying each observations using the set of trees which were built without seeing that observation [2].

The effect of individual features on the classification outcomes can be less transparent for nonlinear classifiers, like random forest, when compared to linear regression-based classifiers. However, there are several ways to recover the variable importance for the classification trees in discriminating between classes. We report the mean variable importance of the random forest classifiers determined by the mean decrease in the Gini index across all branches [2]. The Gini index measures node purity increase, or the increase in the homogeneity of a class in a given node that results from division of the features space based on a particular feature. Informally, for binary classification, variable importance measures how well a feature divides the feature space into partitions that contain a more class-homogeneous subset of the observations. In addition to variable importance we compare the difference of feature distributions for each class of EEG.




Figure 5.1: Spectral features segmented on changes in the B complexity coefficient.

5.5 Model Performance

We begin by defining several terms we will be using to describe classifier performance. We term a seizure a positive response or simply a response. A *true positive* is an accurate prediction of a seizure while a *true negative* is an accurate prediction of a non-response. The *sensitivity* of the classifier is then defined as the proportion of true positives to total positives

$$\frac{\text{True positives}}{\text{Total Positives}}$$

and *specificity* is the total of true negatives to total negatives:

$$\frac{\text{True Negatives}}{\text{Total Negatives}}.$$

Because we are using a small dataset with somewhat imbalanced classes, we report balanced accuracy in place of accuracy, and we refer to this value simply as 'accuracy'.

$$\text{Balanced Accuracy} = \frac{\text{Sensitivity} + \text{Specificity}}{2}.$$

We also report AUC, or area under the ROC curve: see Figure 5.5 for an example ROC curve from two of the classifiers tested. The curve is the plot of the sensitivity against the false positive rate or 1–specificity and a higher AUC corresponds to a better performing classifier. An AUC of 1 indicates the correct classification of all observations.

A baseline classification model was built using the mean of each of the 8 features on the six channel resulting in 48 features. A random forest classifier was trained on these features and the out-of-bag classification results are reported in table 5.1. This baseline model classifies the over represented class, the negative or no-seizure responses, with 93% accuracy, but classification of a seizure response was no better than random.

	Sensitivity	Specificity	Accuracy
1	0.544	0.935	0.740

Table 5.1: Classification performance of baseline classifier.

We compared six models to this baseline model. For each model a different partition scheme was used. For the models using complexity coefficients we label the models A , B , $A + B$, where the label corresponds to the complexity coefficient or coefficients whose change points determined the partition. For example, the $A + B$ model partitions features based on the union of the change points of both complexity

coefficients A and B . For three other models we partitioned features into regular segments of length 8, 15 or 30. We refer to these models by the number of partitions.

The performance of these models was assessed using 5-fold cross-validation where balanced sets were used for the hold-out set for each fold. The reported results are based on 50 bootstrap samples and the confidence intervals give the empirical 95% bootstrap confidence intervals for each metric. In general, all models performed best on the LFP channels, here labeled channel 1 and 2. The best performing model was model 8 followed by the B model.




Figure 5.2: Diagnostic values with 95% bootstrap confidence intervals for each classifier trained on channel 1 features.

The mean and 95% confidence intervals for the performance measures for each classifier using features computed on channel are shown in Figure 5.2. All methods classified non-response trials well but the two better performing models —model 8 and model B —had relatively high accuracy in classifying seizure responses: 76%

for model B and 79% for model 8.

	CH1	(sd1)	CH2	(sd2)	CH3	(sd3)	CH4	(sd4)	CH5	(sd5)	CH6	(sd6)
A	0.82	(0.04)	0.77	(0.04)	0.69	(0.02)	0.71	(0.04)	0.57	(0.04)	0.63	(0.03)
B	0.83	(0.03)	0.77	(0.04)	0.71	(0.03)	0.71	(0.03)	0.61	(0.05)	0.65	(0.06)
A+B	0.81	(0.03)	0.73	(0.03)	0.70	(0.04)	0.72	(0.03)	0.57	(0.09)	0.63	(0.07)
8	0.83	(0.05)	0.82	(0.04)	0.74	(0.04)	0.72	(0.02)	0.75	(0.05)	0.74	(0.06)
15	0.82	(0.05)	0.84	(0.04)	0.73	(0.04)	0.73	(0.02)	0.77	(0.05)	0.78	(0.06)
30	0.79	(0.07)	0.77	(0.05)	0.76	(0.04)	0.76	(0.03)	0.73	(0.08)	0.75	(0.04)

Table 5.2: Mean and standard deviation of balanced accuracy for all models.

	CH1	(sd1)	CH2	(sd2)	CH3	(sd3)	CH4	(sd4)	CH5	(sd5)	CH6	(sd6)
A	0.73	(0.06)	0.60	(0.08)	0.46	(0.04)	0.52	(0.05)	0.26	(0.07)	0.33	(0.05)
B	0.76	(0.07)	0.59	(0.07)	0.51	(0.06)	0.53	(0.05)	0.37	(0.09)	0.38	(0.11)
A+B	0.69	(0.07)	0.51	(0.06)	0.49	(0.08)	0.53	(0.05)	0.23	(0.17)	0.32	(0.12)
8	0.77	(0.11)	0.70	(0.07)	0.59	(0.07)	0.54	(0.04)	0.62	(0.11)	0.57	(0.11)
15	0.71	(0.09)	0.73	(0.08)	0.57	(0.06)	0.56	(0.00)	0.68	(0.12)	0.62	(0.12)
30	0.64	(0.14)	0.60	(0.09)	0.60	(0.08)	0.61	(0.06)	0.59	(0.15)	0.58	(0.09)

Table 5.3: Mean and standard deviation of specificity for all models.

The balanced accuracy of all models for each channel is reported in Table 5.2. All partition schemes resulted in fairly high rates of accuracy for channels 1 and 2. The models with a higher number of partitions performed significantly better on the iEEG channels 3-6. Accurate classification of seizures tailed off sharply for all iEEG channels as seen in Table 5.3. In particular, specificity was greater than 70% for only a few models, all of which used channels 1 or 2. The ROC and smoothed ROC curves in Figure 5.5 for are similar for both the model $A + B$ and model 8.

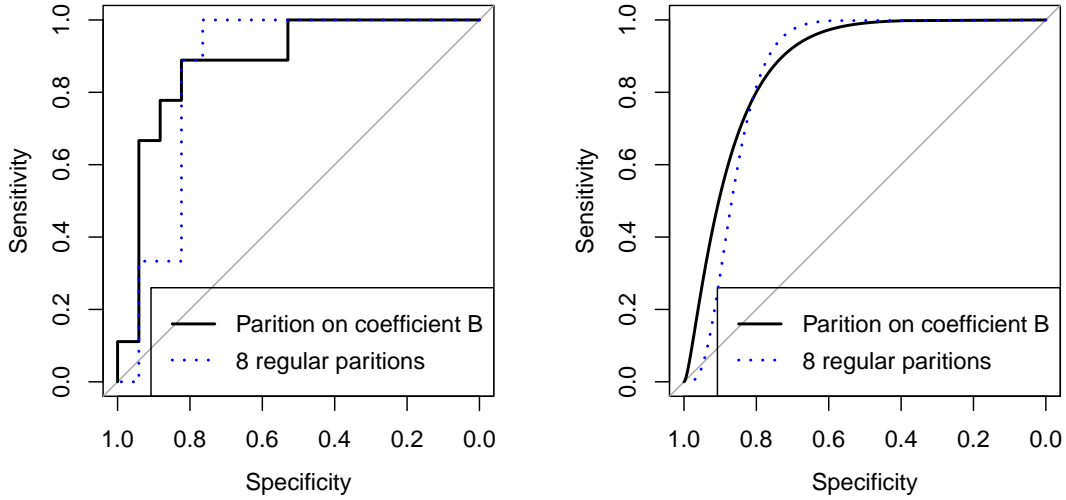


Figure 5.3: Best ROC curves for classifiers B and 8. On the right is the smoothed curve.

5.6 Feature Importance

The classification results show that seizures could be predicted with a relatively high degree of accuracy for some model and channel combinations. The baseline model aside, each of these models used features calculated on individual channels. This allows us to look at the combination of feature and channel associated with accurate predictions. Here we compare the variable importance of the more successful models to the difference in feature distributions between trials with a seizure and non-seizure response. Random forests build multiple decision trees that divide up the feature space in complex ways and the final classification is tallied from the votes of individual trees. This means there may be no clear relation between feature distribution and random forest variable importance. We found, however, that features associated with the best predictors were also associated with clear differences

in distributions between trial classes.




Figure 5.4: Normalized variable importance for each partition method and channel.

In Figure 5.4 we show the variable importance for the two best performing models B and the model 8. We have normalized variable importance to a $[0, 1]$ interval for each model and channel so the figure shows the relative importance of the variable for each model. There is a common pattern in the variable importance across the two models. Gamma and theta power had the greatest importance for the best performing models, those trained on the LFP channels 1 and 2. Both partition methods also show increased importance for beta on channels 4 and 5 and increase the importance of the Hurst coefficient for channels 3 and 4. However, no model was able to accurately predict seizures using these channels in isolation. The maximum specificity among all models trained on these channels was 68% and for the two models depicted in 5.4 this number was 59%.




Figure 5.5: Feature distribution for channels 1(top) and 3(bottom).

Variable importance as measured by the random forest classifiers was also reflected in distributional differences in the features. The box plots in figure 5.5 show the distribution and features for channels 1 and 3. For channel 1, the relative power of the gamma band is higher and the power of the theta band was lower for trials with seizure responses. The median of theta and gamma power fall outside the inter-quartile range and a similar distribution was similar for channel two. For channel three, the relative power of theta and gamma are reversed with gamma power lower and theta power higher. On the other hand, the distribution of beta power for channels 1 and 2 were similar while beta power was significantly lower for channels 3 and 4. Table 5.4 shows the p-value determined by the unpaired Wilcoxon or Mann Whitney test, a nonparametric test for the difference in the location of two distributions. While the large number of data points that even small differences will be statistically significant, the table does highlight the non-significant values.

Channel	1	2	3	4	5	6
Delta	< .0001	< .0001	< .0001	< .0001	< .0001	< .0001
Theta	< .0001	< .0001	< .0001	< .0001	< .0001	< .0001
Alpha	0.022	0.699	0.897	0.891	0.008	0.302
Beta	0.826	0.226	< .0001	< .0001	< .0001	< .0001
Gamma	< .0001	< .0001	< .0001	< .0001	< .0001	< .0001
Variance	< .0001	< .0001	0.687	0.005	< .0001	< .0001
Hurst	< .0001	< .0001	< .0001	< .0001	< .0001	< .0001
Spectral Entropy	< .0001	< .0001	< .0001	< .0001	0.0003	< .0001

Table 5.4: Results of an unpaired Wilcoxon test.

We report the difference of medians in ???. For each feature and channel the feature values were normalized to $(0, 1)$. The table represents the difference of medians of each feach after features ranges were normalized to a $[0, 1]$ interval. The

value is the difference

$$\text{Seizure median} - \text{Non-seizure median}$$

for each feature and channel.

	Channel 1	Channel 2	Channel 3	Channel 4	Channel 5	Channel 6
Delta	0.04	0.03	0.05	0.05	0.03	0.03
Theta	-0.09	-0.09	0.02	0.02	-0.05	-0.06
Alpha	-0.01	0	0	0	-0.01	0
Beta	0	0	-0.04	-0.04	-0.02	-0.01
Gamma	0.05	0.05	-0.05	-0.05	0.05	0.04
Variance	0	0	0	0	0	0
Hurst	0.01	0.01	0.02	0.02	0.01	0.01
Spec. Entropy	0.02	0.02	-0.04	-0.04	0.01	0.01

Table 5.5: Normalized difference of medians (Seizure – Non-seizure) for each feature and trial.

While we were able to predict a seizure result with relatively high accuracy for several segmentation models, there are some limitations to the study. Data from 4 mice were used in the experiments but the trials resulting in seizures came from a single mouse. Data from trials with seizure responses for additional mice would be required to know whether the features associated with seizures might hold more generally. In addition, most of the seizures came from stimuli applied within two trials. Several of the seizures responses, then, occurred after a previous seizure. Therefore, features found predictive of a seizure may be conflated with features that are a result of a seizure. Due to the small number of trials, we did not have a hold-out data set. The lab which produced the dataset is collecting additional data. Testing the predictive models described here against new data would better indicate whether

their predictive power generalize well. On some channels, the segmentation model based on the complexity coefficients performed as well as the models with regular partitions for some channels. However, the regular partition models performed more consistently across channels. For models trained on the iEEG electrodes, channels 3-6, those that used uniformly partitioned features performed better than those partitioned on change points in the complexity coefficients.

Another source of possible variation in the complexity-coefficient segmentation model is the window on which ε -complexity is calculated. The change-point detection algorithm is sensitive to the density at which the complexity-coefficients were computed. Only two variations in the feature window were tested, one at 4 second intervals and one at 2 second intervals. The 4 second window interval resulted in too few samples to adequately test for change points in the complexity coefficients. The feature extraction step is the most computationally expensive, but computing both the complexity coefficients and the other features on a sliding, overlapping, window would possibly allow for a better detection of change points in the diagnostic sequence of the complexity coefficients. On average, the uniform partitions divided the trials into more segments than when the trials were segmented based on the complexity coefficients. It may be that simply increasing the number of segments in comparison to the complexity coefficient segmentation is what improved the consistency of the models with uniform partitions. Tests of this segmentation model on simulated data sets or a wider range of time series would be needed to assess the performance of this method in more general contexts.

Bibliography

- [1] Turkey N Alotaiby, Saleh A Alshebeili, Tariq Alshawhi, Ishtiaq Ahmad, and Fathi E Abd El-Samie, *EEG seizure detection and prediction algorithms: a survey*, EURASIP Journal on Advances in Signal Processing **2014** (2014), no. 1, 183.
- [2] Leo Breiman, *Random forests*, Machine learning **45** (2001), no. 1, 5–32.
- [3] Michael X Cohen, *Where does EEG come from and what does it mean?*, Trends in Neurosciences (2017).
- [4] Germund Dahlquist and Ake Björck, *Numerical methods in scientific computing. Vol. I*, Society for Industrial and Applied Mathematics (SIAM), Philadelphia, PA, 2008. MR 2412832
- [5] Boris Darkhovsky and Alexandra Piryatinska, *Epsilon-complexity of continuous functions*, (2013).
- [6] Kenneth Falconer, *Fractal geometry*, second ed., John Wiley & Sons, Inc., Hoboken, NJ, 2003, Mathematical foundations and applications.
- [7] Jerome Friedman, Trevor Hastie, and Robert Tibshirani, *The elements of statistical learning*, vol. 1, Springer series in statistics New York, 2001.
- [8] Tilmann Gneiting and Martin Schlather, *Stochastic models that separate fractal dimension and the Hurst effect*, SIAM Rev. **46** (2004), no. 2, 269–282.
- [9] Tilmann et al. Gneiting, *Estimators of fractal dimension: assessing the roughness of time series and spatial data*, Statist. Sci. **27** (2012), no. 2, 247–277.
- [10] Brian R. Hunt, *The Hausdorff dimension of graphs of Weierstrass functions*, Proc. Amer. Math. Soc. **126** (1998), no. 3, 791–800.

- [11] Susumu Ito, Ikuo Ogiwara, Kazuyuki Yamada, Hiroyuki Miyamoto, Takao K Hensch, Makiko Osawa, and Kazuhiro Yamakawa, *Mouse with Na V 1.1 haploinsufficiency, a model for Dravet syndrome, exhibits lowered sociability and learning impairment*, Neurobiology of disease **49** (2013), 29–40.
- [12] Nathanael Aff and Alexandra Piryatinska, *ecomplex: A package for computing epsilon-complexity*, 2017.
- [13] Steven Orey, *Gaussian sample functions and the Hausdorff dimension of level crossings*, Z. Wahrscheinlichkeitstheorie und Verw. Gebiete **15** (1970), 249–256. MR 0279882
- [14] Jianqing Fan; Yao Qiwei, *Nonlinear time series*, Springer, New York, 1993.
- [15] Weixiao Shen, *Hausdorff dimension of the graphs of the classical Weierstrass functions*, arXiv preprint arXiv:1505.03986 (2015).
- [16] Wim Sweldens and Peter Schröder, *Building your own wavelets at home*, Wavelets in the Geosciences (2000), 72–107.
- [17] Paul Vitanyi and Ming Li, *An introduction to Kolmogorov complexity and its applications*, Springer-Verlag, New York, 1993.
- [18] Rafał Weron, *Estimating long-range dependence: finite sample properties and confidence intervals*, Phys. A **312** (2002), no. 1-2, 285–299.